

```
In [ ]: !pip install d2l==v1.0.0-alpha1.post0
!pip install pytorch-ignite
!pip install torchviz
!pip install Graphviz
!pip install colormap
```

The following additional libraries are needed to run this notebook. Note that running on Colab is experimental, please report a Github issue if you have any problem.

```
In [ ]: import torch
from torch import nn
from torch.nn import functional as F
from d2l import torch as d2l
from torchvision import transforms
import torchvision
import time
import numpy as np
import matplotlib.pyplot as plt
import graphviz

from functools import partial
```

```
In [ ]: (torch.__version__)
```

```
Out[ ]: '2.0.1+cu118'
```

```
In [ ]: #
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
In [ ]: transform1 = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))], transforms.Pad(4), transforms.RandomErasing(), transforms.RandomHoriz

transform2 = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

batch_size = 128

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform2)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=0)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz

100%|██████████| 170498071/170498071 [00:13<00:00, 13086138.35it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data

Files already downloaded and verified

```
In [ ]: testloader
```

```
Out[ ]: <torch.utils.data.dataloader.DataLoader at 0x7fa1e956aec0>
```

Using a different trainer from d2l

Trainer adjusted

```
In [ ]: def train_ch13(net, trainloader, testloader, num_epochs,
                    device, batchsize):

    net = net.to(device)
    loss = nn.CrossEntropyLoss()
    trainer = torch.optim.SGD(net.parameters(), 0.01)

    timer, num_batches = d2l.Timer(), len(trainloader)
    animator = d2l.Animator(xlabel='epoch', xlim=[1, num_epochs], ylim=[0, 1],
                           legend=['train loss', 'train acc', 'test acc'])
    for epoch in range(num_epochs):
        # Sum of training loss, sum of training accuracy, no. of examples,
        # no. of predictions, LR sched, i
```

```

metric = d2l.Accumulator(4)
for i, (features, labels) in enumerate(trainloader):
    features=features.to(device)
    labels=labels.to(device)
    timer.start()
    l, acc = train_batch(
        net, features, labels, loss, trainer, device)
    metric.add(l, acc, labels.shape[0], labels.numel())
    timer.stop()
    if (i + 1) % (num_batches // 5) == 0 or i == num_batches - 1:
        animator.add(epoch + (i + 1) / num_batches,
            (metric[0] / metric[2], metric[1] / metric[3],
             None))
test_acc = d2l.evaluate_accuracy_gpu(net, testloader)
animator.add(epoch + 1, (None, None, test_acc))

print(f'total test accuracy {test_acc:.3f}')
print(f'Total training time per epoch {timer.sum()/num_epochs:.1f}')
print(f'{ timer.sum():.1f} total training time ')
print(f'{str(device)}')

def train_batch(net,X, y, loss, trainer, device):

    net.train()
    trainer.zero_grad()
    pred = net(X)
    l = loss(pred, y)
    l.sum().backward()
    trainer.step()
    train_loss_sum = l.sum()
    train_acc_sum = d2l.accuracy(pred, y)
    return train_loss_sum, train_acc_sum

```

```

In [ ]: class Residual(nn.Module):
        """The Residual block of ResNet models."""
        def __init__(self, num_channels):
            super().__init__()
            self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                         stride=1)
            self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                         stride=1)
            self.conv3 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                         stride=1)

            self.bn1 = nn.LazyBatchNorm2d()
            self.bn2 = nn.LazyBatchNorm2d()
            self.bn3 = nn.LazyBatchNorm2d()
            self.maxpool = nn.MaxPool2d(kernel_size=2)

            self.ident1 = nn.Identity()
            self.ident2 = nn.Identity()

        def forward(self, X):
            Y = F.relu(self.bn1(self.maxpool(self.conv1(X))),0.075)
            X= self.ident1(Y)
            Y = F.relu(self.bn2(self.conv2(X)),0.075)
            Y = F.relu(self.bn3(self.conv3(Y)),0.075)
            Y= self.ident2(Y)
            Y = X+Y
            return F.relu(Y,0.075)

```

```

In [ ]: class ResNet(d2l.Classifier):

        def __init__(self, arch, lr=0.1, num_classes=10):
            super().__init__()
            self.save_hyperparameters()
            self.net = nn.Sequential(self.b1())

        #Adding blocks
        for i, b in enumerate(arch):
            if i == 1:
                self.net.add_module(f'b{i+2}',nn.Sequential(
                    nn.LazyConv2d(256, kernel_size=3, padding=1,stride=1),
                    nn.MaxPool2d(kernel_size=2),
                    nn.BatchNorm2d(256),
                    nn.CELU(alpha=0.075)
                ))
            else:
                self.net.add_module(f'b{i+2}', self.block(*b, first_block=(i==0)))

        ##Final block
        self.net.add_module('last', nn.Sequential(
            nn.MaxPool2d(4), nn.Flatten(),nn.LazyLinear(num_classes),nn.Identity()))

```

```

self.net.apply(d2l.init_cnn)

def b1(self):
    return nn.Sequential(
        nn.LazyConv2d(64, kernel_size=3, stride=1, padding=1),
        nn.BatchNorm2d(64), nn.CELU(alpha=0.075) )

def block(self, num_residuals, num_channels, first_block=False):
    blk = []
    for i in range(num_residuals):
        blk.append(Residual( num_channels))
    return nn.Sequential(*blk)

def layer_summary(self, X_shape):
    """Defined in :numref:`sec_lenet`"""
    X = d2l.randn(*X_shape)
    for layer in self.net:
        X = layer(X)
        print(layer.__class__.__name__, 'output shape:\t', X.shape)

```

## BaseLine resnet 18

```

In [ ]: transform1 = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),transforms.Pad(4)])

transform2 = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

batch_size = 128

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                       download=True, transform= transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                       shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform2)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                       shuffle=False, num_workers=0)

```

Files already downloaded and verified  
Files already downloaded and verified


Model	Hyper parameters	Accuracy	Time per epoch	Total time
ResNet18	Batchsize:128 Epochs:40 Lr:0.01	67.3	10.2	408.4
	Batchsize:512 Epochs:20	55.1	8.0	159.2

```

In [ ]: net=d2l.resnet18(10,3)
net(next(iter(trainloader))[0])
net=net.to(device)
train_ch13(net, trainloader,testloader, 40,
          device,128)

```

total test accuracy 0.673  
Total training time per epoch 10.2  
408.4 total training time  
cuda:0



Clearly overfitting

## 2ND resnet18

Model	Hyper parameters	Accuracy	Time per epoch	Total time
ResNet18	Batchsize:128 Epochs:40 Lr:0.01	67.3	10.2	408.4
	Batchsize:512 Epochs:20	55.1	8.0	159.2


```
In [ ]: batch_size = 512

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform2)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=0)

net=d2l.resnet18(10,3)
net(next(iter(trainloader))[0])
net=net.to(device)
train_ch13(net, trainloader,testloader, 20,
           device,512)
```

total test accuracy 0.551  
Total training time per epoch 8.0  
159.2 total training time  
cuda:0



Still Clearly overfitting

```
In [ ]: class ResNet9(ResNet):
        def __init__(self, lr=0.1, num_classes=10):
            super().__init__(((1, 128), (1, 256), (1, 512)),
                             lr, num_classes)
        ResNet18().layer_summary((1,3,32,32))

Sequential output shape: torch.Size([1, 64, 32, 32])
Sequential output shape: torch.Size([1, 128, 16, 16])
Sequential output shape: torch.Size([1, 256, 8, 8])
Sequential output shape: torch.Size([1, 512, 4, 4])
Sequential output shape: torch.Size([1, 10])
```

```
In [ ]: def train_batch(net,X, y, loss, trainer, device):
```

```

net.train()
trainer.zero_grad()
pred = net(X)
l = loss(pred, y)
l.sum().backward()
trainer.step()
train_loss_sum = l.sum()
train_acc_sum = d2l.accuracy(pred, y)
return train_loss_sum, train_acc_sum

```

```
In [ ]: device= torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
In [ ]: batch_size = 128
```

```

trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0, pin_memory=True)

testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=0, pin_memory=True)

```


ResNet9	Batchsize:128	68.5	5.6	56.0
	Epochs:10			
	Lr:0.01			

[link text](#)

```
In [ ]: net=ResNet18()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)

train_ch13(net, trainloader, testloader, 10, device, 128)
```

total test accuracy 0.685  
Total training time per epoch 5.6  
56.0 total training time  
cuda:0



```
In [ ]: len(trainloader)
```

```
Out[ ]: 391
```

more epoch, and larger batch size.

Batchsize:512	59.9	4.9	98.6
Epochs:20			
Lr:0.01			

```
In [ ]: batch_size = 512
```

```

trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0)

testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=0)


```

```

net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader, testloader, 20,
           device, 512)

```

total test accuracy 0.599  
Total training time per epoch 4.9  
98.6 total training time  
cuda:0



Just increasing batch size results in large accuracy losses but substantial training time reductions.

Learning rate max	Accuracy	Time per epoch	Total time
0.8	78.8	3.3	66.6
0.4	78.5	3.3	66.5

```
In [ ]: def train_ch13(net, trainloader, testloader, num_epochs,
               device, batchsize, lr):

    trainer = torch.optim.SGD(net.parameters(), lr=lr) #, weight_decay = 0.00001 momentum= 0.9)
    loss = nn.CrossEntropyLoss()

    LRSched = torch.optim.lr_scheduler.OneCycleLR(trainer,
                                                    max_lr=0.8,
                                                    epochs=num_epochs,
                                                    steps_per_epoch=len(trainloader),
                                                    anneal_strategy = "linear")

    timer, num_batches = d2l.Timer(), len(trainloader)
    animator = d2l.Animator(xlabel='epoch', xlim=[1, num_epochs], ylim=[0, 1],
                            legend=['train acc', 'test acc'])
    for epoch in range(num_epochs):
        # Sum of training loss, sum of training accuracy, no. of examples,
        # no. of predictions, LRSched, i

        metric = d2l.Accumulator(3)
        net.train()
        for i, (X, y) in enumerate(trainloader):
            timer.start()
            trainer.zero_grad()
            X, y = X.to(device), y.to(device)
            y_hat = net(X)
            l = loss(y_hat, y)
            with torch.no_grad():
                l.backward()
            metric.add(l * X.shape[0], d2l.accuracy(y_hat, y), X.shape[0])

            trainer.step()
            LRSched.step()
            timer.stop()

            if (i + 1) % (num_batches // 5) == 0 or i == num_batches - 1:
                animator.add(epoch + (i + 1) / num_batches,
                              (metric[1] / metric[2],
                               None))
        test_acc = d2l.evaluate_accuracy_gpu(net, testloader)
        animator.add(epoch + 1, (None, test_acc))

    print(f'total test accuracy {test_acc:.3f}')
    print(f'Total training time per epoch {timer.sum()/num_epochs:.1f}')
    print(f'{timer.sum():.1f} total training time ')
    print(f'{str(device)}')
```

Learning rate max	Accuracy	Time per epoch	Total time
0.8	78.8	3.3	66.6

```
In [ ]: transform1 = transforms.Compose(
    [
        torchvision.transforms.Resize(40),
        torchvision.transforms.RandomResizedCrop(32, scale=(0.64, 1.0),
                                                    ratio=(1.0, 1.0)),
        torchvision.transforms.RandomHorizontalFlip(),
        torchvision.transforms.ToTensor(),

        torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                         [0.2023, 0.1994, 0.2010])
    ])

transform_test = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                     [0.2023, 0.1994, 0.2010])])

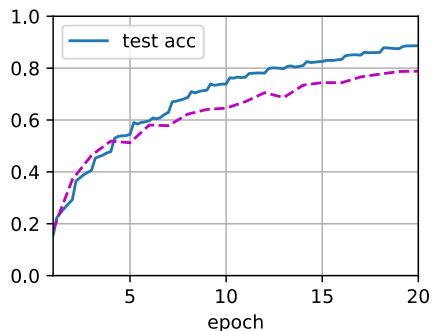
batch_size = 512

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                         shuffle=False, num_workers=0)

net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader, testloader, 20,
           device, 512, 0.01)
```

total test accuracy 0.788  
Total training time per epoch 3.3  
66.6 total training time  
cuda:0



```
In [ ]: def train_ch13(net, trainloader, testloader, num_epochs,
                    device, batchsize, lr):

    trainer = torch.optim.SGD(net.parameters(), lr=lr)#, weight_decay = 0.000001momentum= 0.9)
    loss = nn.CrossEntropyLoss()

    LRSched = torch.optim.lr_scheduler.OneCycleLR(trainer ,
                                                    max_lr=0.4,
                                                    epochs=num_epochs,
                                                    steps_per_epoch=len(trainloader),
                                                    anneal_strategy = "linear")

    timer, num_batches = d2l.Timer(), len( trainloader)
    animator = d2l.Animator(xlabel='epoch', xlim=[1, num_epochs], ylim=[0, 1],
                           legend=['train acc', 'test acc'])
    for epoch in range(num_epochs):
        # Sum of training Loss, sum of training accuracy, no. of examples,
        # no. of predictions, LRSched, i

        metric = d2l.Accumulator(3)
        net.train()
        for i, (X, y) in enumerate( trainloader):
            timer.start()
            trainer.zero_grad()
            X, y = X.to(device), y.to(device)
            y_hat = net(X)
            l = loss(y_hat, y)
            with torch.no_grad():
```

```

        l.backward()
        metric.add(1 * X.shape[0], d2l.accuracy(y_hat, y), X.shape[0])

    trainer.step()
    LRSched.step()
    timer.stop()

    if (i + 1) % (num_batches // 5) == 0 or i == num_batches - 1:
        animator.add(epoch + (i + 1) / num_batches,
                      (metric[1] / metric[2],
                       None))
    test_acc = d2l.evaluate_accuracy_gpu(net, testloader)
    animator.add(epoch + 1, (None, test_acc))

print(f'total test accuracy {test_acc:.3f}')
print(f'Total training time per epoch {timer.sum()/num_epochs:.1f}')
print(f'{ timer.sum():.1f} total training time ')
print(f'{str(device)}')

```

0.4	78.5	3.3	66.5
-----	------	-----	------

```

In [ ]: transform1 = transforms.Compose(
    [
        torchvision.transforms.Resize(40),
        torchvision.transforms.RandomResizedCrop(32, scale=(0.64, 1.0),
                                                  ratio=(1.0, 1.0)),
        torchvision.transforms.RandomHorizontalFlip(),
        torchvision.transforms.ToTensor(),

        torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                         [0.2023, 0.1994, 0.2010])
    ])

transform_test = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                     [0.2023, 0.1994, 0.2010])])

batch_size = 512


trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=0)

net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader, testloader, 20,
           device, 512, 0.01)

```

total test accuracy 0.785  
Total training time per epoch 3.3  
66.5 total training time  
cuda:0



Decay and momentum

Add weight decay and momentum to SGD optimiser. Generalise better and smoothes the SGD line making optimiser more efficient (quicker to convergence). Input linear scheduling. Also removed the plotting of training accuracy and training loss.



```
In [ ]: def train_ch13(net, trainloader, testloader, num_epochs,
                  device, batch_size, lr):

    trainer = torch.optim.SGD(net.parameters(), lr=lr, momentum= 0.9, weight_decay = 0.00001) #
    loss = nn.CrossEntropyLoss()

    LRSched = torch.optim.lr_scheduler.OneCycleLR(trainer,
                                                    max_lr=0.8,
                                                    epochs=num_epochs,
                                                    steps_per_epoch=len(trainloader),
                                                    anneal_strategy = "linear")

    timer, num_batches = d2l.Timer(), len(trainloader)
    animator = d2l.Animator(xlabel='epoch', xlim=[1, num_epochs], ylim=[0, 1],
                            legend=['train acc', 'test acc'])
    for epoch in range(num_epochs):
        # Sum of training loss, sum of training accuracy, no. of examples,
        # no. of predictions, LRSched, i

        metric = d2l.Accumulator(3)
        net.train()
        for i, (X, y) in enumerate(trainloader):
            timer.start()
            trainer.zero_grad()
            X, y = X.to(device), y.to(device)
            y_hat = net(X)
            l = loss(y_hat, y)
            with torch.no_grad():
                l.backward()
                metric.add(1 * X.shape[0], d2l.accuracy(y_hat, y), X.shape[0])

            trainer.step()
            LRSched.step()
            timer.stop()

            if (i + 1) % (num_batches // 5) == 0 or i == num_batches - 1:
                animator.add(epoch + (i + 1) / num_batches,
                              (metric[1] / metric[2],
                               None))

        test_acc = d2l.evaluate_accuracy_gpu(net, testloader)
        animator.add(epoch + 1, (None, test_acc))

    print(f'total test accuracy {test_acc:.3f}')
    print(f'Total training time per epoch {timer.sum()/num_epochs:.1f}')
    print(f'{ timer.sum():.1f} total training time ')
    print(f'{str(device)}')
```

```
In [ ]: transform1 = transforms.Compose(
    [
        torchvision.transforms.Resize(40),
        torchvision.transforms.RandomResizedCrop(32, scale=(0.64, 1.0),
                                                    ratio=(1.0, 1.0)),
        torchvision.transforms.RandomHorizontalFlip(),
        torchvision.transforms.ToTensor(),

        torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                           [0.2023, 0.1994, 0.2010]),
        transforms.RandomErasing())

    transform_test = torchvision.transforms.Compose([
        torchvision.transforms.ToTensor(),
        torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                           [0.2023, 0.1994, 0.2010])])

    batch_size = 512


    trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                              download=True, transform=transform1)
    trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                              shuffle=True, num_workers=0)

    testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                              download=True, transform=transform_test)
    testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                              shuffle=False, num_workers=0)

    net = ResNet9()
    net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
    net = net.to(device)
```

```
train_ch13(net, trainloader, testloader, 20,
           device, 512, 0.01)
```

total test accuracy 0.787  
 Total training time per epoch 3.3  
 66.9 total training time  
 cuda:0



*Epoch and Batch size (81.2%, 133.7s, 40 epochs), (78.8%, 66.9s, 20 epochs) & (67.4%, 64.8s, 20 epochs)*

Varying Epochs from 20-40 at this stage resulted in no substantial gain with only a 2.4% accuracy increase. Furthermore, increasing batch size from 512 to 726 resulted in a 1.2% loss with only a 0.1s gain in speed per epoch. Thus, establishing themselves as not very viable options at this stage.

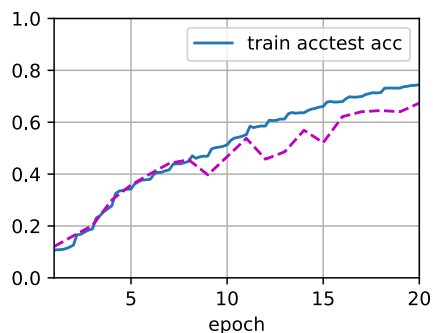
```
In [ ]: batch_size = 726

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                         shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                       shuffle=False, num_workers=0)

net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader, testloader, 20,
           device, 726, 0.01)
```

total test accuracy 0.674  
 Total training time per epoch 3.2  
 64.8 total training time  
 cuda:0




```
In [ ]: batch_size = 512

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                         shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                       shuffle=False, num_workers=0)

net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader, testloader, 20,
           device, 512, 0.01)
```

total test accuracy 0.788  
 Total training time per epoch 3.3  
 66.9 total training time  
 cuda:0



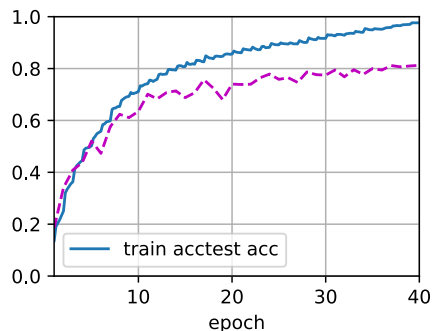
```
In [ ]: batch_size = 512

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                         shuffle=False, num_workers=0)

net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader, testloader, 40,
           device, 512, 0.01)
```

total test accuracy 0.812  
 Total training time per epoch 3.3  
 133.7 total training time  
 cuda:0



Increasing epochs does not seem like a feasible option for increasing accuracy.

```
In [ ]: def train_ch13(net, trainloader, testloader, num_epochs,
                    device, batchsize, lr):

    trainer = torch.optim.SGD(net.parameters(), lr=lr, momentum=0.9, weight_decay=0.00001) #
    loss = nn.CrossEntropyLoss(label_smoothing=0.2)

    LRSched = torch.optim.lr_scheduler.OneCycleLR(trainer,
                                                    max_lr=0.8,
                                                    epochs=num_epochs,
                                                    steps_per_epoch=len(trainloader),
                                                    anneal_strategy="linear")

    timer, num_batches = d2l.Timer(), len(trainloader)
    animator = d2l.Animator(xlabel='epoch', xlim=[1, num_epochs], ylim=[0, 1],
                           legend=['train acc', 'test acc'])
    for epoch in range(num_epochs):
        # Sum of training loss, sum of training accuracy, no. of examples,
        # no. of predictions, LRSched, i

        metric = d2l.Accumulator(3)
        net.train()
        for i, (X, y) in enumerate(trainloader):
            timer.start()
            trainer.zero_grad()
            X, y = X.to(device), y.to(device)
```

```

y_hat = torch.mul(net(X),0.125)
l = loss(y_hat, y)
with torch.no_grad():
    l.backward()
    metric.add(1 * X.shape[0], d2l.accuracy(y_hat, y), X.shape[0])

trainer.step()
LRsched.step()
timer.stop()
if (i + 1) % (num_batches // 5) == 0 or i == num_batches - 1:
    animator.add(epoch + (i + 1) / num_batches,
                  (metric[1] / metric[2],
                   None))

test_acc = d2l.evaluate_accuracy_gpu(net, testloader)
animator.add(epoch + 1, (None, test_acc))

print(f'total test accuracy {test_acc:.3f}')
print(f'Total training time per epoch {timer.sum()/num_epochs:.1f}')
print(f'{ timer.sum():.1f} total training time ')
print(f'{str(device)}')

```

In [ ]:

```

class ResNet(d2l.Classifier):

    def __init__(self, arch, lr=0.1, num_classes=10):
        super().__init__()
        self.save_hyperparameters()
        self.net = nn.Sequential(self.b1())

    #Adding blocks
    for i, b in enumerate(arch):
        if i == 1:
            self.net.add_module(f'b{i+2}',nn.Sequential(
                nn.LazyConv2d(256, kernel_size=3, padding=1,stride=1),
                nn.MaxPool2d(kernel_size=2),
                nn.BatchNorm2d(256,eps=1),
                nn.CELU(alpha=0.075)
            ))
        else:
            self.net.add_module(f'b{i+2}', self.block(*b, first_block=(i==0)))

    ##Final block
    self.net.add_module('last', nn.Sequential(
        nn.MaxPool2d(4), nn.Flatten(),nn.LazyLinear(num_classes),nn.Identity()))
    self.net.apply(d2l.init_cnn)

    def b1(self):
        return nn.Sequential(
            nn.LazyConv2d(64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64,eps=1), nn.CELU(alpha=0.075) )

    def block(self, num_residuals, num_channels, first_block=False):
        blk = []
        for i in range(num_residuals):
            blk.append(Residual( num_channels))
        return nn.Sequential(*blk)

    def layer_summary(self, X_shape):
        """Defined in :numref:`sec_lenet`"""
        X = d2l.randn(*X_shape)
        for layer in self.net:
            X = layer(X)
            print(layer.__class__.__name__, 'output shape:\t', X.shape)

class Residual(nn.Module):
    """The Residual block of ResNet models."""
    def __init__(self, num_channels):
        super().__init__()
        self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                    stride=1)
        self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                    stride=1)
        self.conv3 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                    stride=1)

        self.bn1 = nn.LazyBatchNorm2d(1)
        self.bn2 = nn.LazyBatchNorm2d(1)
        self.bn3 = nn.LazyBatchNorm2d(1)
        self.maxpool = nn.MaxPool2d(kernel_size=2)

```

```

        self.ident1 = nn.Identity()
        self.ident2 = nn.Identity()

    def forward(self, X):
        Y = F.relu(self.bn1(self.maxpool(self.conv1(X))),0.075)
        X= self.ident1(Y)
        Y = F.relu(self.bn2(self.conv2(X)),0.075)
        Y = F.relu(self.bn3(self.conv3(Y)),0.075)
        Y= self.ident2(Y)
        Y = X+Y
        return F.relu(Y,0.075)

```


Label smoothing & Scaling batch norm scales (87.9%, 68.8s, 20 epoch)

```

In [ ]: net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader,testloader, 20,
            device,512,0.01)

```

total test accuracy 0.879  
Total training time per epoch 3.4  
68.8 total training time  
cuda:0



```

In [ ]: def train_ch13(net, trainloader, testloader,num_epochs,
                    device,batchsize,lr):

    trainer=torch.optim.SGD(net.parameters(), lr=lr, momentum= 0.9, weight_decay = 0.00005)#
    loss=nn.CrossEntropyLoss(label_smoothing=0.2)

    LRSched = torch.optim.lr_scheduler.OneCycleLR(trainer ,
                                                    max_lr=0.9,
                                                    epochs=num_epochs,
                                                    steps_per_epoch=len(trainloader),
                                                    anneal_strategy = "linear")

    timer, num_batches = d2l.Timer(), len( trainloader)
    animator = d2l.Animator(xlabel='epoch', xlim=[1, num_epochs], ylim=[0, 1],
                            legend=['train acc','test acc'])

    for epoch in range(num_epochs):
        # Sum of training loss, sum of training accuracy, no. of examples,
        # no. of predictions,LRSched, i

        metric = d2l.Accumulator(3)
        net.train()
        for i, (X, y) in enumerate( trainloader):
            timer.start()
            trainer.zero_grad()
            X, y = X.to(device), y.to(device)
            y_hat = torch.mul(net(X),0.125)
            l = loss(y_hat, y)
            with torch.no_grad():
                l.backward()
                metric.add(l * X.shape[0], d2l.accuracy(y_hat, y), X.shape[0])

            trainer.step()
            LRSched.step()
            timer.stop()
            if (i + 1) % (num_batches // 5) == 0 or i == num_batches - 1:
                animator.add(epoch + (i + 1) / num_batches,
                              (metric[1] / metric[2],
                               None))

        test_acc = d2l.evaluate_accuracy_gpu(net, testloader)
        animator.add(epoch + 1, ( None, test_acc))

```

```

print(f'total test accuracy {test_acc:.3f}')
print(f'Total training time per epoch {timer.sum()/num_epochs:.1f}')
print(f'{ timer.sum():.1f} total training time ')
print(f'{str(device)}')

```

*Epoch and Batch size (90.2%, 191.1s, 50 epochs)*

In [ ]:

```

transform1 = transforms.Compose(
    [
        torchvision.transforms.Resize(40),
        torchvision.transforms.RandomResizedCrop(32, scale=(0.64, 1.0),
                                                    ratio=(1.0, 1.0)),
        torchvision.transforms.RandomHorizontalFlip(),
        torchvision.transforms.ToTensor(),

        torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                         [0.2023, 0.1994, 0.2010]),
        transforms.RandomErasing()])

transform_test = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                     [0.2023, 0.1994, 0.2010])])

batch_size = 256

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0)


testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=0)

device= torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
class ResNet18(ResNet):
    def __init__(self, lr=0.1, num_classes=10):
        super().__init__(((1, 128), (1, 256), (1, 512)),
                        lr, num_classes)
ResNet18().layer_summary((1,3,32,32))

net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader, testloader, 50,
           device, 256, 0.01)

```

total test accuracy 0.902  
Total training time per epoch 3.8  
191.1 total training time  
cuda:0



Architecture

In [ ]:

```

class ResNet(d2l.Classifier):

    def __init__(self, arch, lr=0.1, num_classes=10):
        super().__init__()
        self.save_hyperparameters()
        self.net = nn.Sequential(self.b1())

    #Adding blocks
    for i, b in enumerate(arch):
        if i == 1:
            self.net.add_module(f'b{i+2}', nn.Sequential(
                nn.LazyConv2d(256, kernel_size=3, padding=1, stride=1),

```

```

        nn.BatchNorm2d(256,eps=1),
        nn.MaxPool2d(kernel_size=2),
        nn.CELU(alpha=0.075)
    ))
    else:
        self.net.add_module(f'b{i+2}', self.block(*b, first_block=(i==0)))

##Final block
self.net.add_module('last', nn.Sequential(
    nn.MaxPool2d(4), nn.Flatten(),nn.LazyLinear(num_classes),nn.Identity()))
self.net.apply(d2l.init_cnn)

def b1(self):
    return nn.Sequential(
        nn.LazyConv2d(64, kernel_size=3, stride=1, padding=1),
        nn.BatchNorm2d(64,eps=1), nn.CELU(alpha=0.075) )

def block(self, num_residuals, num_channels, first_block=False):
    blk = []
    for i in range(num_residuals):
        blk.append(Residual( num_channels))
    return nn.Sequential(*blk)

def layer_summary(self, X_shape):
    """Defined in :numref:`sec_lenet`"""
    X = d2l.randn(*X_shape)
    for layer in self.net:
        X = layer(X)
        print(layer.__class__.__name__, 'output shape:\t', X.shape)

class Residual(nn.Module):
    """The Residual block of ResNet models."""
    def __init__(self, num_channels):
        super().__init__()
        self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                    stride=1)
        self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                    stride=1)
        self.conv3 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                    stride=1)

        self.bn1 = nn.LazyBatchNorm2d(1)
        self.bn2 = nn.LazyBatchNorm2d(1)
        self.bn3 = nn.LazyBatchNorm2d(1)
        self.maxpool = nn.MaxPool2d(kernel_size=2)

        self.ident1 = nn.Identity()
        self.ident2 = nn.Identity()

    def forward(self, X):
        Y = F.celu(self.maxpool(self.bn1(self.conv1(X))),0.075)
        X= self.ident1(Y)
        Y = F.celu(self.bn2(self.conv2(X)),0.075)
        Y = F.celu(self.bn3(self.conv3(Y)),0.075)
        Y= self.ident2(Y)
        Y = X+Y
        return F.celu(Y,0.075)

```

```

In [ ]: transform1 = transforms.Compose(
    [

        torchvision.transforms.Resize(40),
        torchvision.transforms.RandomResizedCrop(32, scale=(0.64, 1.0),
                                                    ratio=(1.0, 1.0)),
        torchvision.transforms.RandomHorizontalFlip(),
        torchvision.transforms.ToTensor(),

        torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                         [0.2023, 0.1994, 0.2010]),
        transforms.RandomErasing())

transform_test = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                     [0.2023, 0.1994, 0.2010])])

batch_size = 512

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                         shuffle=True, num_workers=0)


testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform_test)

```

```
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                         shuffle=False, num_workers=0)

device= torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader,testloader, 20,
           device,512,0.01)
```

total test accuracy 0.871  
Total training time per epoch 3.7  
73.0 total training time  
cuda:0



Drop layers

```
In [ ]: class ResNet(d2l.Classifier):

    def __init__(self, arch, lr=0.1, num_classes=10):
        super().__init__()
        self.save_hyperparameters()
        self.net = nn.Sequential(self.b1())

    #Adding blocks
    for i, b in enumerate(arch):
        if i == 1:
            self.net.add_module(f'b{i+2}',nn.Sequential(
                nn.Dropout(p=0.2),
                nn.LazyConv2d(256, kernel_size=3, padding=1,stride=1),
                nn.MaxPool2d(kernel_size=2),
                nn.BatchNorm2d(256,eps=1),
                nn.CELU(alpha=0.075)
            ))
        else:
            self.net.add_module(f'b{i+2}', self.block(*b, first_block=(i==0)))

    ##Final block
    self.net.add_module('last', nn.Sequential(nn.Dropout(0.7),
        nn.MaxPool2d(4), nn.Flatten(),nn.LazyLinear(num_classes),nn.Identity()))
    self.net.apply(d2l.init_cnn)

    def b1(self):
        return nn.Sequential(
            nn.LazyConv2d(64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(64,eps=1), nn.CELU(alpha=0.075) )

    def block(self, num_residuals, num_channels, first_block=False):
        blk = []
        for i in range(num_residuals):
            blk.append(Residual( num_channels))
        return nn.Sequential(*blk)

    def layer_summary(self, X_shape):
        """Defined in :numref:`sec_lenet`"""
        X = d2l.randn(*X_shape)
        for layer in self.net:
            X = layer(X)
            print(layer.__class__.__name__, 'output shape:\t', X.shape)

class Residual(nn.Module):
    """The Residual block of ResNet models."""
    def __init__(self, num_channels):
        super().__init__()
        self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                   stride=1)
        self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
```



```

        stride=1)
self.conv3 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                           stride=1)

self.bn1 = nn.LazyBatchNorm2d(1)
self.bn2 = nn.LazyBatchNorm2d(1)
self.bn3 = nn.LazyBatchNorm2d(1)
self.maxpool = nn.MaxPool2d(kernel_size=2)
self.Dropout=nn.Dropout(0.2)
self.ident1 = nn.Identity()
self.ident2 = nn.Identity()

def forward(self, X):
    Y= self.Dropout(X)
    Y = F.relu(self.bn1(self.maxpool(self.conv1(X))),0.075)
    X= self.ident1(Y)
    Y = F.relu(self.bn2(self.conv2(X)),0.075)
    Y = F.relu(self.bn3(self.conv3(Y)),0.075)
    Y= self.ident2(Y)
    Y = X+Y
    return F.relu(Y,0.075)

```

### Architecture

Finally, the models had their maxpool and batch norms switched. This was done as it was recommended and resulting in a reduction in accuracy.

```

In [ ]: transform1 = transforms.Compose(
    [
        torchvision.transforms.Resize(40),
        torchvision.transforms.RandomResizedCrop(32, scale=(0.64, 1.0),
                                                ratio=(1.0, 1.0)),
        torchvision.transforms.RandomHorizontalFlip(),
        torchvision.transforms.ToTensor(),

        torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                         [0.2023, 0.1994, 0.2010]),
        transforms.RandomErasing())

transform_test = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize([0.4914, 0.4822, 0.4465],
                                     [0.2023, 0.1994, 0.2010])])

batch_size = 512

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform1)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                           shuffle=True, num_workers=0)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                          shuffle=False, num_workers=0)

device= torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
net=ResNet9()
net.apply_init([next(iter(trainloader))[0]], d2l.init_cnn)
net=net.to(device)
train_ch13(net, trainloader,testloader, 20,
           device,512,0.01)

```

total test accuracy 0.870  
 Total training time per epoch 3.5  
 69.4 total training time  
 cuda:0

