

The following additional libraries are needed to run this notebook. Note that running on Colab is experimental, please report a Github issue if you have any problem.

```
In [ ]: !pip install d2l==v1.0.0-alpha1.post0
!pip install pytorch-ignite
!pip install torchviz
```

```
In [2]: import torch
from torch import nn
from torch.nn import functional as F
from d2l import torch as d2l
from torchvision import transforms
import torchvision
import time
import numpy as np
```

```
In [3]: device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
```

```
In [4]: class FashionMNIST(d2l.DataModule):
    """The Fashion-MNIST dataset.

    Defined in :numref:`sec_fashion_mnist`"""
    def __init__(self, batch_size=64, resize=(28, 28)):
        super().__init__()
        self.save_hyperparameters()
        trans = transforms.Compose([transforms.Resize(resize), transforms.ToTensor()])

        self.train = torchvision.datasets.FashionMNIST(
            root=self.root, train=True, transform=trans, download=True)
        self.val = torchvision.datasets.FashionMNIST(
            root=self.root, train=False, transform=trans, download=True)
    def text_labels(self, indices):
        """Return text labels.

        Defined in :numref:`sec_fashion_mnist`"""
        labels = ['t-shirt', 'trouser', 'pullover', 'dress', 'coat',
                  'sandal', 'shirt', 'sneaker', 'bag', 'ankle boot']
        return [labels[int(i)] for i in indices]

    def get_dataloader(self, train):
        """Defined in :numref:`sec_fashion_mnist`"""
        data = self.train if train else self.val
        return torch.utils.data.DataLoader(data, self.batch_size, shuffle=train,
                                            num_workers=self.num_workers)

    def visualize(self, batch, nrows=1, ncols=8, labels=[]):
        """Defined in :numref:`sec_fashion_mnist`"""
        X, y = batch
        if not labels:
            labels = self.text_labels(y)
        d2l.show_images(X.squeeze(1), nrows, ncols, titles=labels)
```

```
In [5]: class Trainer(d2l.HyperParameters):
    def __init__(self, max_epochs, num_gpus=0, gradient_clip_val=0):
        """Defined in :numref:`sec_use_gpu`"""
        self.save_hyperparameters()
        self.gpus = [d2l.gpu(i) for i in range(min(num_gpus, d2l.num_gpus()))]
        self.avg_accuracy = []
        self.avgtimer=None
        self.epochtime = None

    def prepare_data(self, data):
        self.train_dataloader = data.train_dataloader()
        self.val_dataloader = data.val_dataloader()
        self.num_train_batches = len(self.train_dataloader)
        self.num_val_batches = (len(self.val_dataloader)
                                if self.val_dataloader is not None else 0)

    def fit(self, model, data):
        timer=d2l.Timer()
        self.prepare_data(data)
        self.prepare_model(model)
        self.optim = model.configure_optimizers()
        self.sched = torch.optim.lr_scheduler.OneCycleLR(self.optim,
                                                         max_lr=0.4,
                                                         epochs=self.max_epochs,
                                                         steps_per_epoch=len(self.train_dataloader),
                                                         pct_start= 0.25,
                                                         anneal_strategy = "linear")

        self.epoch = 0
        self.train_batch_idx = 0
```

```

self.val_batch_idx = 0
for self.epoch in range(self.max_epochs):
    timer.start()
    self.fit_epoch()
    timer.stop()
    self.model.epoch_accuracy()
    self.epochtime=timer

def fit_epoch(self):
    raise NotImplementedError

def prepare_batch(self, batch):
    """Defined in :numref:`sec_linear_scratch`"""
    return batch

def fit_epoch(self):
    """Defined in :numref:`sec_linear_scratch`"""
    self.model.train()
    for batch in self.train_dataloader:
        loss = self.model.training_step(self.prepare_batch(batch))
        self.optim.zero_grad()
        with torch.no_grad():
            loss.backward()
            if self.gradient_clip_val > 0: # To be discussed later
                self.clip_gradients(self.gradient_clip_val, self.model)
            self.optim.step()
            self.sched.step()
        self.train_batch_idx += 1
    if self.val_dataloader is None:
        return
    self.model.eval()
    for batch in self.val_dataloader:
        with torch.no_grad():
            accuracy=self.model.validation_step(self.prepare_batch(batch))

    self.val_batch_idx += 1

def prepare_batch(self, batch):
    """Defined in :numref:`sec_use_gpu`"""
    if self.gpus:
        batch = [d2l.to(a, self.gpus[0]) for a in batch]
    return batch

def prepare_model(self, model):
    """Defined in :numref:`sec_use_gpu`"""
    model.trainer = self
    model.board.xlim = [0, self.max_epochs]
    if self.gpus:
        model.to(self.gpus[0])
    self.model = model

def clip_gradients(self, grad_clip_val, model):
    """Defined in :numref:`sec_rnn_scratch`"""
    params = [p for p in model.parameters() if p.requires_grad]
    norm = torch.sqrt(sum(torch.sum((p.grad ** 2)) for p in params))
    if norm > grad_clip_val:
        for param in params:
            param.grad[:] *= grad_clip_val / norm

```

Global maxpooling before the output layer

```

In [13]: class varied_Residual(nn.Module):
    """The Residual block of ResNet models."""
    def __init__(self, num_channels, use_1x1conv=False, strides=1):
        super().__init__()
        #self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
        #                           stride=strides)
        #self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1)

        self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3,
                                    stride=1)

        self.bn1 = nn.LazyBatchNorm2d()
        self.maxpool = nn.MaxPool2d(kernel_size=2)

    def forward(self, X):
        #Y = F.relu(self.bn1(self.conv1(X)))
        Y = F.relu(self.bn1(self.conv1(X)))
        Y=self.maxpool(Y)
        return Y

```

In [14]:

```
class ResNet(d2l.Classifier):

    def __init__(self, arch, lr=0.1, num_classes=10):
        super().__init__()
        self.save_hyperparameters()
        self.net = nn.Sequential(self.b1())
        self.averageaccuracy = []
        self.epoch_accuracy_vals = []

#Adding blocks
    for i, b in enumerate(arch):
        self.net.add_module(f'b{i+2}', self.block(*b, first_block=(i==0)))

##Final block

    self.net.add_module('last', nn.Sequential(
        nn.MaxPool2d(4), nn.Flatten(),
        nn.Linear(num_classes)))
    self.net.apply(d2l.init_cnn)

    def b1(self):
        return nn.Sequential(
            nn.Conv2d(64, kernel_size=3, stride=1, padding=1),
            nn.BatchNorm2d(), nn.ReLU())

    def block(self, num_residuals, num_channels, first_block=False):
        blk = []
        for i in range(num_residuals):
            if i == 0 and not first_block:
                blk.append(varied_Residual(num_channels, use_1x1conv=True, strides=1))
            else:
                blk.append(varied_Residual(num_channels))
        return nn.Sequential(*blk)

#####
    def epoch_accuracy(self):
        epoch_acc = torch.mean(torch.stack(self.averageaccuracy))
        self.averageaccuracy = []
        self.epoch_accuracy_vals.append(epoch_acc)

    def validation_step(self, batch):
        Y_hat = self(*batch[:-1])
        self.plot('loss', self.loss(Y_hat, batch[-1]), train=False)
        self.plot('acc', self.accuracy(Y_hat, batch[-1]), train=False)
        accuracy = self.accuracy(Y_hat, batch[-1])
        self.averageaccuracy.append(accuracy)

    def accuracy(self, Y_hat, Y, averaged=True):
        """Compute the number of correct predictions.

        Defined in :numref:`sec_classification`"""
        Y_hat = d2l.reshape(Y_hat, (-1, Y_hat.shape[-1]))
        preds = d2l.astype(d2l.argmax(Y_hat, axis=1), Y.dtype)
        compare = d2l.astype(preds == d2l.reshape(Y, -1), d2l.float32)
        return d2l.reduce_mean(compare) if averaged else compare

    def loss(self, Y_hat, Y, averaged=True):
        """Defined in :numref:`sec_softmax_concise`"""
        Y_hat = d2l.reshape(Y_hat, (-1, Y_hat.shape[-1]))
        Y = d2l.reshape(Y, (-1,))
        return F.cross_entropy(
            Y_hat, Y, reduction='mean' if averaged else 'none')

    def layer_summary(self, X_shape):
        """Defined in :numref:`sec_lenet`"""
        X = d2l.randn(*X_shape)
        for layer in self.net:
            X = layer(X)
            print(layer.__class__.__name__, 'output shape:\t', X.shape)
```

In [15]:

```
class ResNet9(ResNet):
    def __init__(self, lr=0.1, num_classes=10):
        super().__init__((1, 128), (1, 256), (1, 512)),
            lr, num_classes)

ResNet18()
ResNet18().layer_summary((1,1,96,96))
```

```

Sequential output shape: torch.Size([1, 64, 96, 96])
Sequential output shape: torch.Size([1, 128, 47, 47])
Sequential output shape: torch.Size([1, 256, 22, 22])
Sequential output shape: torch.Size([1, 512, 10, 10])
Sequential output shape: torch.Size([1, 10])

```

In [16]:

```

model = ResNet9(lr=0.04)

data = FashionMNIST(batch_size=512, resize=(96, 96))

trainer = Trainer(max_epochs=4, num_gpus=1)
model.apply_init([next(iter(data.get_dataloader(True)))[0]], d2l.init_cnn)

trainer.fit(model, data)

for x in range(len(model.epoch_accuracy_vals)):
    print(f'Epoch {x+1}')
    print(f'Test accuracy: {float(model.epoch_accuracy_vals[x]*100)}%')
    print(f'Epoch time: {trainer.epochtime.times[x]} s')



print(f'Test accuracy: {float(model.epoch_accuracy_vals[len(model.epoch_accuracy_vals)-1]*100)}%')
print(f'Average time: {trainer.epochtime.avg()} s')
print(f'Total time: {trainer.epochtime.sum()} s')

```

```

Epoch 1
Test accuracy: 74.5289535522461
Epoch time: 23.869215965270996 s
Epoch 2
Test accuracy: 79.73287963867188
Epoch time: 23.9197735786438 s
Epoch 3
Test accuracy: 82.34202575683594
Epoch time: 24.023970365524292 s
Epoch 4
Test accuracy: 82.8670654296875
Epoch time: 24.04045271873474 s
Test accuracy: 82.8670654296875
Average time: 23.963353157043457 s
Total time: 95.85341262817383 s

```



Epoch 1 Test accuracy: 74.5289535522461 Epoch time: 23.869215965270996 s Epoch 2 Test accuracy: 79.73287963867188 Epoch time: 23.9197735786438 s Epoch 3 Test accuracy: 82.34202575683594 Epoch time: 24.023970365524292 s Epoch 4 Test accuracy: 82.8670654296875 Epoch time: 24.04045271873474 s Test accuracy: 82.8670654296875 Average time: 23.963353157043457 s Total time: 95.85341262817383 s

#Conversion of ReLU activation functions to CELU and reordering of maxpooling and batch norms

In [17]:

```

class Residual(nn.Module):
    """The Residual block of ResNet models."""
    def __init__(self, num_channels):
        super().__init__()
        self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                   stride=1)
        self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                   stride=1)
        self.conv3 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,

```

```

        stride=1)

self.bn1 = nn.LazyBatchNorm2d()
self.bn2 = nn.LazyBatchNorm2d()
self.bn3 = nn.LazyBatchNorm2d()
self.maxpool = nn.MaxPool2d(kernel_size=2)

def forward(self, X):
    X = nn.CELU(self.bn1(self.maxpool(self.conv1(X))), alpha=0.075)

    Y = nn.CELU(self.bn2(self.conv2(X)), alpha=0.075)
    Y = nn.CELU(self.bn3(self.conv3(Y)), alpha=0.075)
    Y += X
    return F.relu(Y)

```

In [18]:

```

class ResNet(d2l.Classifier):

    def __init__(self, arch, lr=0.1, num_classes=10):
        super().__init__()
        self.save_hyperparameters()
        self.net = nn.Sequential(self.b1())
        self.averageaccuracy = []
        self.epoch_accuracy_vals = []

    #Adding blocks
    for i, b in enumerate(arch):
        if i == 1:
            self.net.add_module(f'b{i+2}', nn.Sequential(
                nn.LazyConv2d(256, kernel_size=3, padding=1, stride=1),
                nn.MaxPool2d(kernel_size=2),
                nn.LazyBatchNorm2d(),
                nn.CELU(alpha=0.075)
            ))

            self.net.add_module(f'b{i+2}', self.block(*b, first_block=(i==0)))

    ##Final block
    self.net.add_module('last', nn.Sequential(
        nn.MaxPool2d(4), nn.Flatten(), nn.LazyLinear(num_classes)))
    self.net.apply(d2l.init_cnn)

    def b1(self):
        return nn.Sequential(
            nn.LazyConv2d(64, kernel_size=3, stride=1, padding=1),
            nn.LazyBatchNorm2d(), nn.ReLU())

    def block(self, num_residuals, num_channels, first_block=False):
        blk = []
        for i in range(num_residuals):
            if i == 0 and not first_block:
                blk.append(varied_Residual(num_channels, use_1x1conv=True, strides=1))
            else:
                blk.append(varied_Residual(num_channels))
        return nn.Sequential(*blk)

    #####
    def epoch_accuracy(self):
        epoch_acc = torch.mean(torch.stack(self.averageaccuracy))
        self.averageaccuracy = []
        self.epoch_accuracy_vals.append(epoch_acc)

    def validation_step(self, batch):
        Y_hat = self(*batch[:-1])
        self.plot('loss', self.loss(Y_hat, batch[-1]), train=False)
        self.plot('acc', self.accuracy(Y_hat, batch[-1]), train=False)
        accuracy = self.accuracy(Y_hat, batch[-1])
        self.averageaccuracy.append(accuracy)

    def accuracy(self, Y_hat, Y, averaged=True):
        """Compute the number of correct predictions.

        Defined in :numref:`sec_classification`"""
        Y_hat = d2l.reshape(Y_hat, (-1, Y_hat.shape[-1]))
        preds = d2l.astype(d2l.argmax(Y_hat, axis=1), Y.dtype)
        compare = d2l.astype(preds == d2l.reshape(Y, -1), d2l.float32)
        return d2l.reduce_mean(compare) if averaged else compare

    def loss(self, Y_hat, Y, averaged=True):
        """Defined in :numref:`sec_softmax_concise`"""
        Y_hat = d2l.reshape(Y_hat, (-1, Y_hat.shape[-1]))
        Y = d2l.reshape(Y, (-1,))

```

```

        return F.cross_entropy(
            Y_hat, Y, reduction='mean' if averaged else 'none')

    def layer_summary(self, X_shape):
        """Defined in :numref:`sec_lenet`"""
        X = d2l.randn(*X_shape)
        for layer in self.net:
            X = layer(X)
            print(layer.__class__.__name__, 'output shape:\t', X.shape)

```

In [19]:

```

class ResNet9(ResNet):
    def __init__(self, lr=0.1, num_classes=10):
        super().__init__((1, 128), (1, 256), (1, 512)),
                        lr, num_classes)

model = ResNet9(lr=0.04)

data = FashionMNIST(batch_size=512, resize=(96, 96))

trainer = Trainer(max_epochs=4, num_gpus=1)
model.apply_init([next(iter(data.get_dataloader(True)))[0]], d2l.init_cnn)

trainer.fit(model, data)

for x in range(len(model.epoch_accuracy_vals)):
    print(f'Epoch {x+1}')
    print(f'Test accuracy: {float(model.epoch_accuracy_vals[x]*100)}')
    print(f'Epoch time: {trainer.epochtime.times[x]} s')



print(f'Test accuracy: {float(model.epoch_accuracy_vals[len(model.epoch_accuracy_vals)-1]*100)}')
print(f'Average time: {trainer.epochtime.avg()} s')
print(f'Total time: {trainer.epochtime.sum()} s')

```

```

Epoch 1
Test accuracy: 70.29756927490234
Epoch time: 23.873865604400635 s
Epoch 2
Test accuracy: 77.23345947265625
Epoch time: 23.99230980873108 s
Epoch 3
Test accuracy: 80.84846496582031
Epoch time: 23.99388575553894 s
Epoch 4
Test accuracy: 83.31629180908203
Epoch time: 24.057832956314087 s
Test accuracy: 83.31629180908203
Average time: 23.979473531246185 s
Total time: 95.91789412498474 s

```



```

Epoch 1 Test accuracy: 70.29756927490234 Epoch time: 23.873865604400635 s Epoch 2 Test accuracy: 77.23345947265625 Epoch time:
23.99230980873108 s Epoch 3 Test accuracy: 80.84846496582031 Epoch time: 23.99388575553894 s Epoch 4 Test accuracy: 83.31629180908203
Epoch time: 24.057832956314087 s Test accuracy: 83.31629180908203 Average time: 23.979473531246185 s Total time: 95.91789412498474 s

```

In []:

